

# EXPERIMENT-1

## WEEK -1: MATRIX OPERATIONS

**AIM: To perform various matrix operations.**

- a) Create multi-dimensional arrays and find its shape and dimension
- b) Create a matrix full of zeros and ones
- c) Reshape and flatten data in the array
- d) Append data vertically and horizontally
- e) Apply indexing and slicing on array
- f) Use statistical functions on array-Min, Max, Mean, Median and Standard

Deviation

### RESOURCES:

- a) Python 3.7.0
- b) Install: pip installer, Pandas library

### PROCEDURE:

1. Create: Open a new file in Python shell, write a program and save the program with .py extension.
2. Execute: Goto Run->Run module(F5)

### PROGRAM LOGIC:

- a) Create multi-dimensional arrays and find its shape and dimension

```
Import numpy as np
#creationofmulti-dimensionalarray
a=np.array ([[1,2,3],[2,3,4],[3,4,5]])

#shape

b=a.shape
print("shape:",a.shape)
```

```
#dimensio
n c=a.ndim
print("dimensions:",a.ndim)
```

**OUTPUT:**

```
shape:(3,3)
dimensions:2
```

b) Create a matrix full of zeros and ones

```
#matrixfullofzeros
```

```
z=np.zeros((2,2))
```

```
print("zeros:",z)
```

```
#matrixfullofones
```

```
o=np.ones((2,2))
```

```
print("ones:",o)
```

**OUTPUT:**

```
zeros:[[0.0.]
```

```
      a. 0.]]
```

```
      ones:[
```

```
      [1. 1.]
```

```
      b. 1.]]
```

c) Reshape and flatten data in the array

### **#matrixreshape**

```
a=np.array([[1,2,3,4],[2,3,4,5],[3,4,5,6],[4,5,6,7]])  
b=a.reshape(4,2,2)  
print("reshape:",b)
```

### **#matrix**

```
flattenc=a.flatten()  
print("flatten: ",c)
```

### **OUTPUT:**

```
reshape:[[[12]  
[34]]  
  
[[23]  
[45]]  
  
[[34]  
[56]]  
  
[[45]  
[6 7]]]  
flatten:[1234234534564567]
```

Regd. No

d) Append data vertically and horizontally

### **#Appending data vertically**

```
x=np.array([[10,20],[80,90]])  
y=np.array([[30,40],[60,70]])  
v=np.vstack((x,y))  
print("vertically:",v)
```

### **#Appending data horizontally**

```
h=np.hstack((x,y))
```

```
)  
    print("horizontally:",h
```

**OUTPUT:**

```
vertically:[[1020]  
[80 90]  
[3040] [6070]]  
horizontally:[[10203040]  
[809060 70]]
```

e) Apply indexing and slicing on array

## **#indexing**

```
a=np.array([[1,2,3,4],[  
2,3,4,5],[3,4,5,6],[4,5,6,7]])  
t  
emp=a[[0,1,2,3],[1,1,1,1]]
```

```
print('indexing',temp)
#slicing i=a[:4,::2] print('slicing',i)
```

**OUTPUT:**

```
indexing[2 3 4 5]
```

```
slicing[[1 3]
```

```
[2 4]
```

```
[3 5]
```

```
[4 6]]
```

f) Use statistical functions on array-Min, Max, Mean, Median and Standard Deviation

**#minforfindingminimumofanarray**

```
a=np.array([[1,3,-1,4],[3,-2,1,4]])
```

```
b=a.min()
```

```
print('minimum'
```

```
',b)
```

**#maxforfindingmaximumofanarray**

```
c=a.max() print('maximum',c)
```

```
a=np.array([1,2,3,4,5])
```

```
d=a.mean ()
```

```
print('mean:',d)
```

**#median**

```
e=np.median( a)
    print('median:',e)

    #standarddeviation
f=a.std()
```

```
print('standarddeviation',f)
```

**OUTPUT:**

minimum:-2

maximum 4

mean:3.0

median:3.0

standarddeviation1.4142135623730951



## **EXPERIMENT -5**

### **WEEK - 5: DATA PREPROCESSING-HANDLING MISSING VALUES**

**Write a python program to input missing values with various techniques on given dataset.**

- a) Remove rows/attributes
- b) Replace with mean or mode
- c) Write a python program to perform transformation of data using Discretization (Binning) and normalization(Min Max Scale or Max AbsScaler) on given dataset.

```
import pandas as pd  
from sklearn.impute import SimpleImputer  
from sklearn.preprocessing import KBinsDiscretizer, MinMaxScaler, MaxAbsScaler
```

**# Load your dataset**

```
data = {  
    'A': [1, 2, None, 4, 5],  
    'B': [None, 2, 3, None, 5],  
    'C': [1, 2, 3, 4, 5]  
}  
df = pd.DataFrame(data)
```

**# Display the original dataframe**

```
print("Original DataFrame:") print(df)
```

**# Handling missing values**

```
def handle_missing_values(df, method='mean'): if method ==  
    'remove':
```

---

Regd. No

df = df.dropna()

**# Remove rows with missing values else:**

if method == 'mean':

imputer = SimpleImputer(strategy='mean') elif method

== 'mode': imputer =

SimpleImputer(strategy='most\_frequent') df[df.columns]

= imputer.fit\_transform(df[df.columns])

return df

**# Apply missing value handling function**

**# You can change the method parameter to 'remove', 'mean', or 'mode'**

df\_handled = handle\_missing\_values(df, method='mean')

**# Display the dataframe after handling missing values**

print("\nDataFrame After Handling Missing Values:")

print(df\_handled)

**# Data Transformation** def data\_transformation(df,

method='binning'): if method == 'binning':

est = KBinsDiscretizer(n\_bins=3, encode='ordinal', strategy='uniform')

df[df.columns] = est.fit\_transform(df[df.columns])

elif method == 'normalization': scaler

= MinMaxScaler()

df[df.columns]

# You can also use MaxAbsScaler for MaxAbs

normalization scaler.fit\_transform(df[df.columns]) return df

**# Apply data transformation function**

**# You can change the method parameter to 'binning' or 'normalization'**

```
df_transformed = data_transformation(df_handled, method='normalization')
```

**# Display the dataframe after data transformation**

```
print("\nDataFrame      After Data Transformation:")
```

```
print(df_transformed)
```

### **OUTPUT:**

Original DataFrame:

	A	B	C
0	1.0	NaN	1
1	2.0	2.0	2
2	NaN	3.0	3
3	4.0	NaN	4
4	5.0	5.0	5

DataFrame After Handling Missing Values: A

	B	C
0	1.0 3.333333	1.0 1
1	2.0 2.000000	2.0
2	3.0 3.000000	3.0
3	4.0 3.333333	4.0
4	5.0 5.000000	5.0

DataFrame After Data Transformation: A

	B	C
0	0.444444	0.0
0.00	4	0
1	0.000000	0.2
0.25	0	5
2	0.333333	0.5
0.50	3	0
3	0.444444	0.7
0.75	4	5
4	1.000000	1.0
1.00	0	0

[ ]

## EXPERIMENT – 7

### WEEK -7: CLASSIFICATION-LOGISTIC REGRESSION

#### LOGISTIC REGRESSION

A statistical model for binary classification is called [logistic regression](#). Using the sigmoid function, it forecasts the likelihood that an instance will belong to a particular class, guaranteeing results between 0 and 1. To minimize the log loss, the model computes a linear combination of input characteristics, transforms it using the sigmoid, and then optimizes its coefficients using methods like gradient descent. These coefficients establish the decision boundary that divides the classes. Because of its ease of use, interpretability, and versatility across multiple domains, Logistic Regression is widely used in machine learning for problems that involve binary outcomes. Overfitting can be avoided by implementing regularization.

#### SOURCE CODE :

##### # Import necessary libraries

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

import seaborn as sns
from sklearn.datasets import load_diabetes
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix, roc_curve, auc

```

##### # Load the diabetes dataset

```

diabetes = load_diabetes()
X, y = diabetes.data, diabetes.target

```

##### # Convert the target variable to binary (1 for diabetes, 0 for no diabetes) y\_binary

```

y_binary = (y > np.median(y)).astype(int)

```

##### # Split the data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y_binary,
test_size=0.2, random_state=42)
```

### # Standardize features

```
scaler = StandardScaler()
```

```
X_train = scaler.fit_transform(X_train) X_test = scaler.transform(X_test)
```

### # Train the Logistic Regression model

```
model = LogisticRegression() model.fit(X_train, y_train)
```

### # Evaluate the model

```
y_pred = model.predict(X_test) accuracy =
accuracy_score(y_test, y_pred) print("Accuracy:
{:.2f}%".format(accuracy * 100))
```

### # evaluate the model

```
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred)) #
```

### Visualize the decision boundary with accuracy information

```
plt.figure(figsize=(8, 6))sns.scatterplot(x=X_test[:, 2], y=X_test[:, 8],
hue=y_test, palette={ 0: 'blue', 1: 'red'}, marker='o')

plt.xlabel("BMI")

plt.ylabel("Age")

plt.title("Logistic Regression Decision Boundary\nAccuracy: {:.2f}%".format( accuracy
* 100))

plt.legend(title="Diabetes", loc="upper right")
plt.show()
```

### # Split the data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(
X, y_binary, test_size=0.2, random_state=42)
```

**# Standardize features**

```
scaler = StandardScaler()
```

```
X_train = scaler.fit_transform(X_train) X_test = scaler.transform(X_test)
```

**# Train the Logistic Regression model**

```
model = LogisticRegression() model.fit(X_train,
y_train)
```

**# Evaluate the model**

```
y_pred = model.predict(X_test) accuracy =
accuracy_score(y_test, y_pred) print("Accuracy:
{:.2f}%".format(accuracy * 100))
```

**# evaluate the model**

```
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

**# Visualize the decision boundary with accuracy information**

```
plt.figure(figsize=(8, 6))
```

```
sns.scatterplot(x=X_test[:, 2], y=X_test[:, 8],
hue=y_test, palette={0: 'blue', 1: 'red'}, marker='o')
plt.xlabel("BMI") plt.ylabel("Age")
```

```
plt.title("Logistic Regression Decision Boundary\nAccuracy: {:.2f}%".format( accuracy *
100))plt.legend(title="Diabetes", loc="upper right") plt.show()
```

Accuracy: 73.03% Confusion  
Matrix:[[36 13][11 29]]

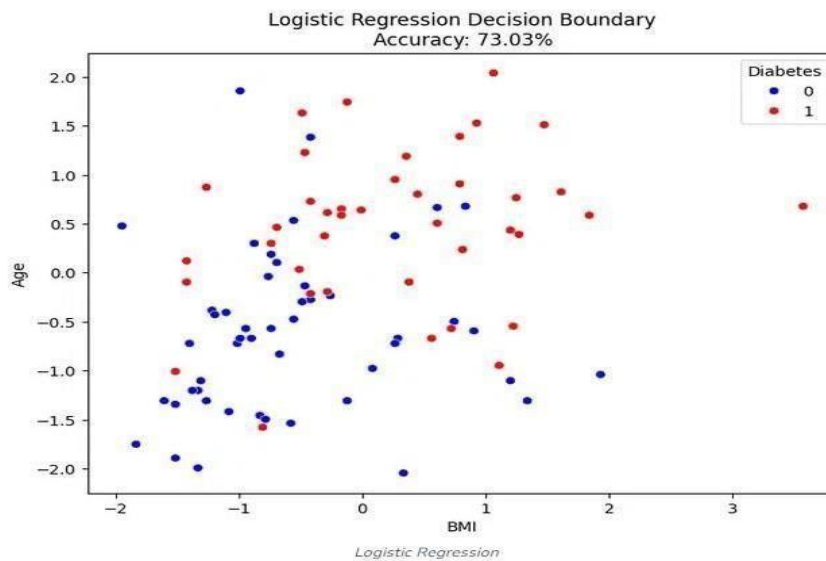
**Classification Report:**

precision recall f1-score support

0	0.77	0.73	0.75	49
0.69	0.72	0.71		40

89 8  
 0.73 macroavg 0.73  
 0.73 0.73 9  
 0.7  
 weighteda 0.7 0.73 3  
 vg 3  
 89

Output:



## EXPERIMENT – 8

### WEEK -8 : CLASSIFICATION-KNN ALGORITHM

#### K-NEAREST NEIGHBOR ALGORITHM:

This algorithm is used to solve the classification model problems. K-nearest neighbor or K-NN algorithm basically creates an imaginary boundary to classify the data. When new data points come in, the algorithm will try to predict that to the nearest of the boundary line.

Therefore, larger k value means smoother curves of separation resulting in less complex models. Whereas, smaller k value tends to overfit the data and result in complex models.

**Note:** It's very important to have the right k-value when analyzing the dataset to avoid overfitting and algorithm we fit the historical data (or train the model) and predict the future.

Here in the example shown above, we are creating a plot to see the k-value for which we have high accuracy.

**Note:** This is a technique which is not used industry-wide to choose the correct value of n-neighbors. Instead, we do hyperparameter tuning to choose the value that gives the best performance. We will be covering this in future posts

#### Source code:

# Import necessary modules from sklearn.neighbors



```
import kNeighborsClassifier from sklearn.model_selection
import train_test_split from sklearn.datasets import load_iris
import numpy as np
import matplotlib.pyplot as plt
irisData = load_iris()
```

### # Create feature and target arrays

```
X = irisData.data
y = irisData.target
```

### # Split into training and test set

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=42)
```

```
neighbors = np.arange(1, 9)
train_accuracy = np.empty(len(neighbors))
test_accuracy = np.empty(len(neighbors))
```

### # Loop over K values for i, k in enumerate(neighbors):

```
knn = KNeighborsClassifier(n_neighbors=k)
knn.fit(X_train, y_train)
```

### # Compute training and test data accuracy

```
train_accuracy[i] = knn.score(X_train, y_train)
test_accuracy[i] = knn.score(X_test, y_test)
```

### # Generate plot

```
plt.plot(neighbors, test_accuracy, label = 'Testing dataset Accuracy')
plt.plot(neighbors, train_accuracy, label = 'Training dataset Accuracy')
plt.legend()
plt.xlabel('n_neighbors')
plt.ylabel('Accuracy')
plt.show()
```

**OUTPUT:**

